

Deploying models with AWS Lambda

- This tutorial is based on <https://github.com/alexeygrigorev/serverless-deep-learning>
- We will deploy a model for predicting the types of clothes (trained here: <https://github.com/alexeygrigorev/mlbookcamp-code/blob/master/chapter-07-neural-nets/07-neural-nets-train.ipynb>)
- You can find all the code here: <https://github.com/alexeygrigorev/aws-lambda-model-deployment-workshop>

Plan:

- Create the needed resources in AWS (optional)
- Convert the model from Keras to TF Lite
- Extract all the pre-processing logic
- Prepare the code for lambda
- Package the dependencies along with the code into a single archive

Prerequisites

- You need to have Python 3.7 (or Python 3.8). The easiest way to install it — use Anaconda (<https://www.anaconda.com/products/individual>)
- Install TensorFlow (pip install tensorflow should be sufficient)
- You need to have an account in AWS

AWS Preparation work

First, we need to do some prep work:

- Create a bucket for storing the code of our lambda function and the model
- Create the lambda function
- Make sure the lambda function can read from the bucket

Log in to AWS console

Create an S3 bucket — we will use it for storing the model and the code of the lambda function. Go to Services ⇒ S3. Click “Create bucket”. Write a name (“lambda-model-deployment-workshop”). For this workshop, we’ll use the same bucket, so you can skip this step.

Create bucket

Buckets are containers for data stored in S3. [Learn more](#) 

General configuration

Bucket name

Bucket name must be unique and must not contain spaces or uppercase letters. [See rules for bucket naming](#) 

Region

EU (Ireland) eu-west-1 ▼

Copy settings from existing bucket - *optional*

Only the bucket settings in the following configuration are copied.

[Choose bucket](#)

Press “create bucket” (at the end)

Create a lambda function. Go to services, select “Lambda”. Click “Create function”.

Select “create from scratch”, and fill in the basic information: function name, runtime (Python 3.7). The rest can stay as is.

Click “create function”.

Basic information

Function name
Enter a name that describes the purpose of your function.

alexey-clothes-classification

Use only letters, numbers, hyphens, or underscores with no spaces.

Runtime [Info](#)
Choose the language to use to write your function.

Python 3.7

Permissions [Info](#)
By default, Lambda will create an execution role with permissions to upload logs to Amazon CloudWatch Logs. You can customize this default role later when adding triggers.

[▶ Change default execution role](#)

[▶ Advanced settings](#)

Cancel

Create function

Now we have a function! We need to make sure this function can read from the S3 bucket we just created — it will load the model from there.

Go to the “permissions” tab, click on the role name to edit it

alexey-clothes-classification

[Configuration](#)[Permissions](#)[Monitoring](#)

Execution role

Role name

alexey-clothes-classification-role-l568eb4p

Select the policy, click on “edit policy”

[Roles](#) > alexey-clothes-classification-role-l568eb4p

Summary

Delete role

Role ARN	arn:aws:iam::[redacted]:role/service-role/alexey-clothes-classification-role-l568eb4p
Role description	Edit
Instance Profile ARNs	
Path	/service-role/
Creation time	2020-11-11 11:47 UTC+0100
Last activity	Not accessed in the tracking period
Maximum session duration	1 hour Edit

Permissions

Trust relationships

Tags

Access Advisor

Revoke sessions

▼ Permissions policies (1 policy applied)

Attach policies

[+ Add inline policy](#)

Policy name ▼	Policy type ▼	
▼ AWSLambdaBasicExecutionRole-979b8c...	Managed policy	✕
<div>Policy summary {} JSON Edit policy Simulate policy</div> <pre>1 { 2 "Version": "2012-10-17", 3 "Statement": [</pre>		

Select the tab with “JSON” and add the following statement:

```
{
  "Effect": "Allow",
  "Action": [
    "s3:Get*"
  ],
  "Resource": [
    "arn:aws:s3:::lambda-model-deployment-workshop",
    "arn:aws:s3:::lambda-model-deployment-workshop/*"
  ]
}
```

Where “lambda-model-deployment-workshop” is the name of the bucket we just created — replace it if your bucket is different.

The full policy should look similar to that:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "logs:CreateLogGroup",
      "Resource": "arn:aws:logs:eu-west-1:XXXXXXXXXXXX:*"
    },
    {
      "Effect": "Allow",
      "Action": [
        "logs:CreateLogStream",
        "logs:PutLogEvents"
      ],
      "Resource": [
        "arn:aws:logs:eu-west-1:XXXXXXXXXXXX:log-group:/aws/lambda/alexey-clothes-classification:*"
      ]
    },
    {
      "Effect": "Allow",
      "Action": [
        "s3:Get*"
      ],
      "Resource": [
        "arn:aws:s3:::lambda-model-deployment-workshop",
        "arn:aws:s3:::lambda-model-deployment-workshop/*"
      ]
    }
  ]
}
```

Then review it and save it.

Preparing the model

Suppose we already trained a model using Keras. Now we want to serve it with AWS Lambda. We need to do a few things for that:

- Convert the model to TF-lite format
- Upload the result to the S3 bucket

We'll do that in a Jupyter notebook (add link)

Get the model:

```
wget
https://github.com/alexeygrigorev/mlbookcamp-code/releases/download/chapter7-
model/xception_v4_large_08_0.894.h5
```

Open a Jupyter notebook (or create a simple python script). Start with the imports:

```
import numpy as np
import tensorflow as tf
from tensorflow import keras
```

Load the model:

```
model = keras.models.load_model('xception_v4_large_08_0.894.h5')
```

Convert it to TF-lite:

```
converter = tf.lite.TFLiteConverter.from_keras_model(model)

tflite_model = converter.convert()

with tf.io.gfile.GFile('clothing-model-v4.tflite', 'wb') as f:
    f.write(tflite_model)
```

The model is ready and we can upload it to S3. Put it to
s3://lambda-model-deployment-workshop/clothing-model-v4.tflite:

```
aws s3 cp clothing-model-v4.tflite
s3://lambda-model-deployment-workshop/clothing-model-v4.tflite
```

Preprocessing functions

To apply the model, we need to do the following steps:

- Get the image (as a PIL Image)
- Prepare the image (resize, etc)
- Convert the image to a tensor, apply the pre-processing function (normalization, etc)
- Put the tensor in the model, get the predictions and post-process the predictions

In Keras, the logic for doing most of these operations is in the `keras-preprocessing` module. We can't use this module inside AWS Lambda (it's too heavy), so we need to write this code ourselves.

Let's do it! For reference, check the notebook [here](#). Later, we'll put this code to our lambda function.

```
from io import BytesIO
from urllib import request

import numpy as np
from PIL import Image

def download_image(url):
    with request.urlopen(url) as resp:
        buffer = resp.read()
    stream = BytesIO(buffer)
    img = Image.open(stream)
    return img

def prepare_image(img, target_size=(224, 224)):
    if img.mode != 'RGB':
        img = img.convert('RGB')
    img = img.resize(target_size, Image.NEAREST)
    return img

def image_to_array(img):
    return np.array(img, dtype='float32')

def tf_preprocessing(x):
    x /= 127.5
    x -= 1.0
    return x

def convert_to_tensor(img):
    x = image_to_array(img)
```

```

batch = np.expand_dims(x, axis=0)
return tf_preprocessing(batch)

```

Note: for some models (resnet, vgg), we need to use caffe preprocessing instead of tf preprocessing:

```

mean = [103.939, 116.779, 123.68]

```

```

def caffe_preprocessing(x):
    # 'RGB' -> 'BGR'
    x = x[..., ::-1]

    x[..., 0] -= mean[0]
    x[..., 1] -= mean[1]
    x[..., 2] -= mean[2]

    return x

```

This is how we can use this code to get a tensor:

```

img = download_image(url)
img = prepare_image(img, target_size=(299, 299))
X = convert_to_tensor(img)

```

Now let's use this code in a model!

Loading the model

Load the model:

- Download it from s3
- Load the actual model from disk

Downloading the model is easy: we just use boto3 for that:

```

import os
import boto3

s3_client = boto3.client('s3')

model_bucket = 'lambda-model-deployment-workshop'
model_key = 'clothing-model-v4.tflite'
model_local_path = '/tmp/clothing-model-v4.tflite'

```



```
if not os.path.exists(model_local_path):
    s3_client.download_file(model_bucket, model_key, model_local_path)
```

To use the model, we first need to load it with TF lite:

```
# import tensorflow.lite as tflite # if testing locally
import tflite_runtime.interpreter as tflite

interpreter = tflite.Interpreter(model_path=model_local_path)
interpreter.allocate_tensors()

input_details = interpreter.get_input_details()
input_index = input_details[0]['index']

output_details = interpreter.get_output_details()
output_index = output_details[0]['index']
```

Now we can use it:

```
interpreter.set_tensor(input_index, X)
interpreter.invoke()

preds = interpreter.get_tensor(output_index)
```

The preds array contains the predictions

Code for Lambda

Each lambda function should have an entrypoint. Let's create it:

```
def lambda_handler(event, context):
    img = download_image(event['url'])
    pred = predict(img)
    result = decode_predictions(pred)
    return result
```

The predict function is just the code from the previous sections put together

```
def predict(img):
```

```

img = prepare_image(img, target_size=(299, 299))
X = convert_to_tensor(img)

interpreter.set_tensor(input_index, X)
interpreter.invoke()

preds = interpreter.get_tensor(output_index)

return preds[0]

```

The `decode_prediction` function turn the raw output into the final result:

```

labels = [
    'dress',
    'hat',
    'longsleeve',
    'outwear',
    'pants',
    'shirt',
    'shoes',
    'shorts',
    'skirt',
    't-shirt'
]

def decode_predictions(pred):
    result = {c: float(p) for c, p in zip(labels, pred)}
    return result

```

We put all this code in `lambda_function.py` (see [the full example](#)).

Preparing the package

To deploy something to AWS Lambda, we need to prepare a zip archive that contains everything: the code itself and all the dependencies.

AWS Lambda uses Amazon Linux, so we need to make sure we use the proper binaries when installing the dependencies. The best way to do this is to package everything in Docker.

So, let's create a Dockerfile for that. We can name it `build.dockerfile`:

```
FROM amazonlinux
```

```
WORKDIR /tflite
```

```
RUN yum groupinstall -y development
```

```
RUN yum install -y python3.7
```

```
RUN yum install -y python3-devel
```

```
RUN pip3 install wheel
```

```
WORKDIR /app
```

```
COPY tflite_runtime-2.2.0-cp37-cp37m-linux_x86_64.whl
```

```
tflite_runtime-2.2.0-cp37-cp37m-linux_x86_64.whl
```

```
# here we don't use requirements.txt/Pipenv for simplicity
```

```
RUN pip3 install \  
    numpy==1.16.5 \  
    Pillow==6.2.1 \  
    tflite_runtime-2.2.0-cp37-cp37m-linux_x86_64.whl \  
    -t build
```

```
WORKDIR /app/build
```

```
RUN (find . -name "tests" | xargs -n1 rm -rf) && \  
    (find . -name "*.pyc" | xargs -n1 rm -rf) && \  
    (find . -name "__pycache__" | xargs -n1 rm -rf) && \  
    (find . -name "*.dist-info" | xargs -n1 rm -rf)
```

```
COPY lambda_function.py lambda_function.py
```

```
RUN zip -r ../build.zip * > /dev/null
```

```
WORKDIR /app
```

```
ENTRYPOINT [ "cp", "build.zip", "results/build.zip" ]
```

For that you'll need a version of TF-Lite compiled for AWS Lambda:

```
wget
```

```
https://github.com/alexeygrigorev/serverless-deep-learning/raw/master/tflite/  
tflite_runtime-2.2.0-cp37-cp37m-linux_x86_64.whl
```

(Use instructions from <https://github.com/alexeygrigorev/serverless-deep-learning> to compile it yourself for other versions of Python)

Let's build an image:

```
BUILDER_IMAGE_NAME=tfllite_build_lambda  
docker build -t ${BUILDER_IMAGE_NAME} -f build.dockerfile .
```

And get the archive out of it:

```
docker run --rm \  
    -v $(pwd):/app/results \  
    ${BUILDER_IMAGE_NAME}
```

After executing it, you'll have a zip archive. You can use it now to upload it to AWS Lambda

Testing the package

Next, we need to check that the package can be unpacked correctly and no dependencies are missing.

Let's create a test.docker file with the following content:

```
FROM amazonlinux  
  
RUN yum groupinstall -y development  
RUN yum install -y python3.7 python3-devel  
RUN pip3 install boto3  
  
WORKDIR /app  
  
COPY build.zip .  
RUN unzip build.zip  
  
COPY test.py test.py  
  
ENV AWS_DEFAULT_REGION="local"  
ENV AWS_ACCESS_KEY_ID="KEY"  
ENV AWS_SECRET_ACCESS_KEY="SECRET"  
  
ENTRYPOINT [ "python3", "test.py" ]
```

Now, create test.py:

```
import lambda_function
```

```

event = {
    "url":
"https://raw.githubusercontent.com/alexeygrigorev/clothing-dataset-small/master/test/pants/4aabd82c-82e1-4181-a84d-d0c6e550d26d.jpg"
}

print(lambda_function.handler(event, None))

```

We're ready to build the image:

```

TEST_IMAGE_NAME=tflite_test_lambda
docker build -t ${TEST_IMAGE_NAME} -f test.dockerfile .

```

And run it to test that our flow works correctly

```

docker run --rm \
    -v $(pwd)/clothing-model-v4.tflite:/tmp/clothing-model-v4.tflite \
    ${TEST_IMAGE_NAME}

```

We should get:

```

{'dress': -1.8682900667190552, 'hat': -4.7612457275390625, 'longsleeve':
-2.3169822692871094, 'outwear': -1.062570571899414, 'pants':
9.88715648651123, 'shirt': -2.8124303817749023, 'shoes': -3.66628360748291,
'shorts': 3.2003610134124756, 'skirt': -2.6023387908935547, 't-shirt':
-4.835044860839844}

```

So our model is working!

Updating the Lambda function

For that we'll need two things:

- Upload the package to S3
- Update the actual lambda function

Let's do it:

```

ZIP_FILE="build.zip"
S3_BUCKET="lambda-model-deployment-workshop"
S3_KEY="alexey/clothes-classification-package.zip"
FUNCTION_NAME="alexey-clothes-classification"

aws s3 cp "${ZIP_FILE}" "s3://${S3_BUCKET}/${S3_KEY}"

```

```
aws lambda update-function-code \  
  --function-name ${FUNCTION_NAME} \  
  --s3-bucket ${S3_BUCKET} \  
  --s3-key ${S3_KEY}
```

Testing the lambda function

Go to the Lambda function.

First, adjust the basic settings. Click edit:

Basic settings [Info](#) Edit

Description	Runtime
-	Python 3.7
Handler Info	Memory (MB)
lambda_function.lambda_handler	128
Timeout	
0 min 3 sec	

Give it 512MB or 1024MB of RAM and set timeout to 30 sec:

Handler [Info](#)

lambda_function.lambda_handler

Memory (MB)

Your function is allocated CPU proportional to the memory configured.

1024 MB

Timeout

min sec

Save it.

Next, create a test with this request:

```
{
  "url":
  "https://raw.githubusercontent.com/alexeygrigorev/clothing-dataset-small/master/test/pants/4aabd82c-82e1-4181-a84d-d0c6e550d26d.jpg"
}
```

Configure test event



A function can have up to 10 test events. The events are persisted so you can switch to another computer or web browser and test your function with the same events.

- ☒ Create new test event
- ☐ Edit saved test events

Event template

hello-world



Event name

Pants

```
1 {
2   "url": "https://raw.githubusercontent.com/alexeygrigorev/clothing-dataset-small/master/test/pants/4aa
3 }
4
```

Save and test it: click the “test” button.

You should see “Execution results: succeeded”:

✓ Execution result: succeeded (logs)

▼ Details

The area below shows the result returned by your function execution. [Learn more](#) about returning results from your function.

```
{
  "dress": -1.8682900667190552,
  "hat": -4.7612457275390625,
  "longsleeve": -2.3169822692871094,
  "outwear": -1.062570571899414,
  "pants": 9.88715648651123,
  "shirt": -2.8124303817749023,
  "shoes": -3.66628360748291,
  "shorts": 3.2003610134124756,
  "skirt": -2.6023387908935547,
  "sleeveless": -2.6023387908935547,
  "sweater": -2.6023387908935547,
  "trousers": -2.6023387908935547,
  "vest": -2.6023387908935547
}
```

Summary

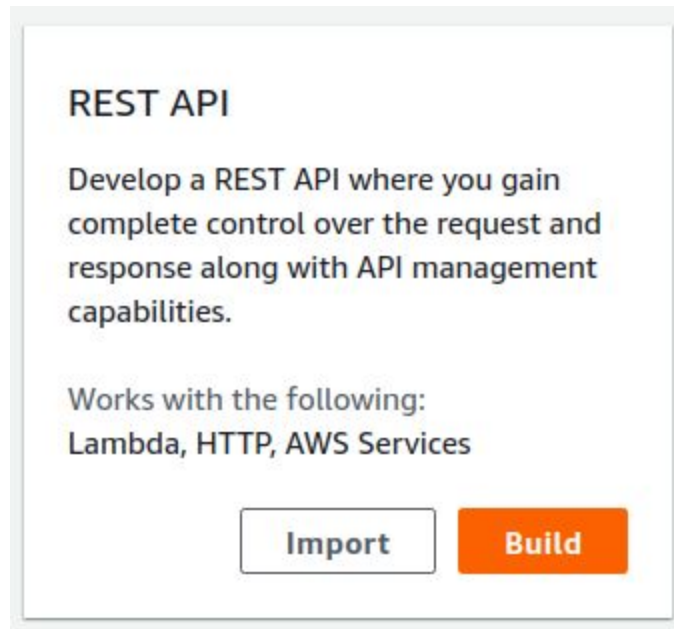
Code SHA-256	Request ID
x2HQV0Zn3CKVdIrfIDPqh1sWb1JvEKl+FFmQ3X08FUE=	ac484545-5a03-491d-ab69-dfd340ddf339
Duration	Billed duration
2907.52 ms	3000 ms
Resources configured	Max memory used
1024 MB	404 MB Init Duration: 2028.72 ms

To be able to use it from outside, we need to create an API. We do it with API Gateway.

Creating the API Gateway

Go to services ⇒ API Gateway

Create a new HTTP API:



Call it "alexey-image-classification"


Create new API

In Amazon API Gateway, a REST API refers to a collection of resources and methods that can be invoked through HTTPS endpoints.

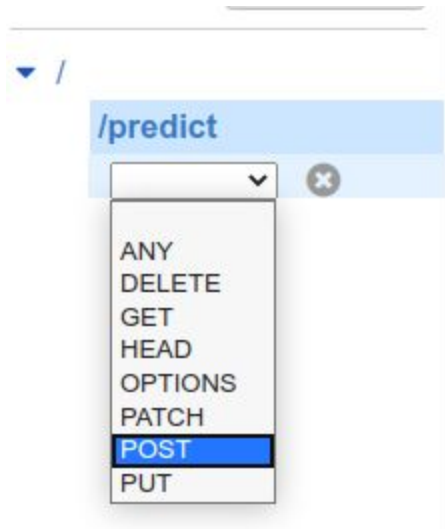
☒ **New API** ☐ Import from Swagger or Open API 3 ☐ Example API

Settings

Choose a friendly name and description for your API.

API name*	<input type="text" value="alexey-image-classification"/>
Description	<input type="text" value="Invoking the lambda function for classification"/>
Endpoint Type	<div>Regional </div>

Then, create a resource "predict", and create a method POST in this resource:



Select “Lambda” and enter the details of your lambda function:

/predict - POST - Setup

Choose the integration point for your new method.

Integration type ☒ Lambda Function ⓘ
☐ HTTP ⓘ
☐ Mock ⓘ
☐ AWS Service ⓘ
☐ VPC Link ⓘ

Use Lambda Proxy integration ☐ ⓘ

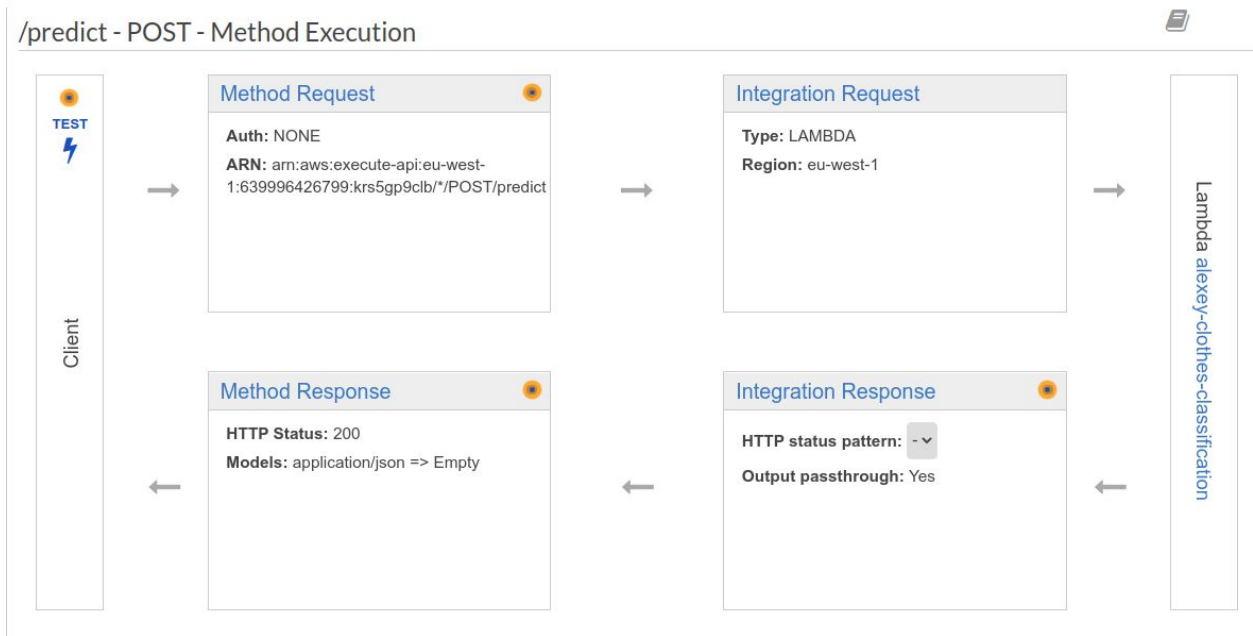
Lambda Region eu-west-1 ▼

Lambda Function alexey-clothes-classification

Use Default Timeout ☒ ⓘ

Make sure you don’t select “proxy integration” — this box should remain unchecked.

Now you should see the integration:



To test it, click on “test” and put this request to request body:

```
{
  "url":
  "https://raw.githubusercontent.com/alexeygrigorev/clothing-dataset-small/master/test/pants/4aabd82c-82e1-4181-a84d-d0c6e550d26d.jpg"
}
```

You should get the response:

Request: /predict

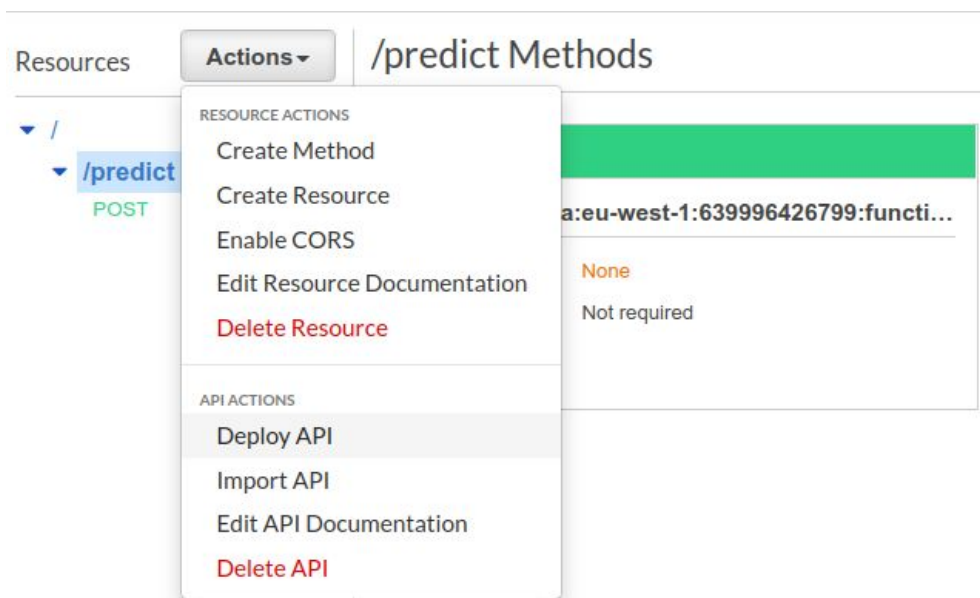
Status: 200

Latency: 1949 ms

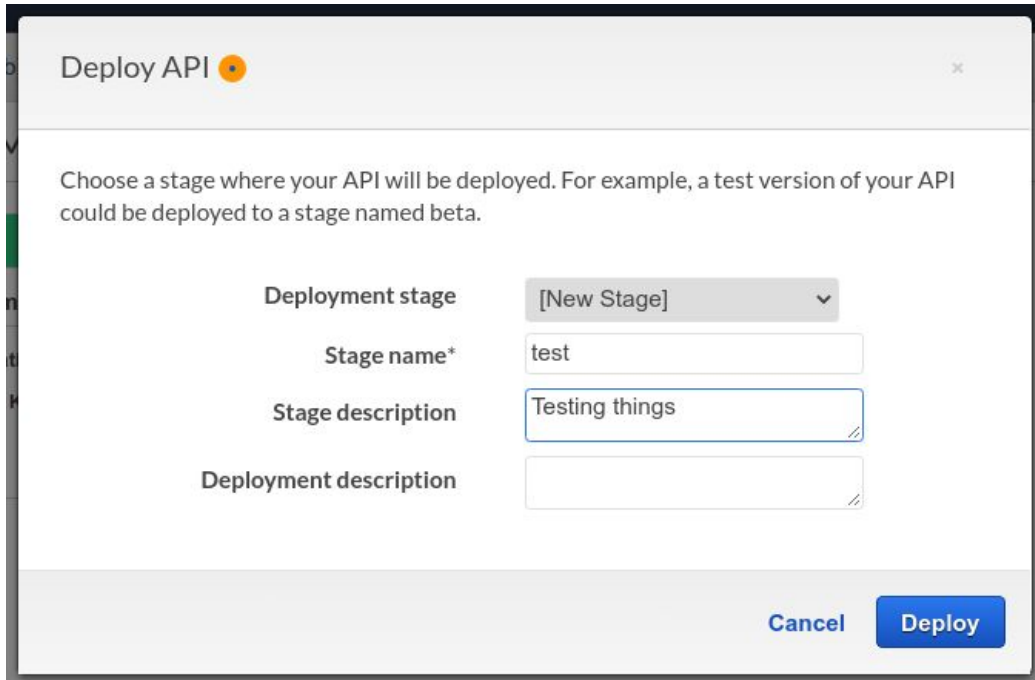
Response Body

```
{
  "dress": -1.8682900667190552,
  "hat": -4.7612457275390625,
  "longsleeve": -2.3169822692871094,
  "outwear": -1.062570571899414,
  "pants": 9.88715648651123,
  "shirt": -2.8124303817749023,
  "shoes": -3.66628360748291,
  "shorts": 3.2003610134124756,
  "skirt": -2.6023387908935547,
  "t-shirt": -4.835044860839844
}
```

To use it, we need to deploy the API. Click on “Deploy API” from Actions.



Create a new stage “test”:



Deploy API

Choose a stage where your API will be deployed. For example, a test version of your API could be deployed to a stage named beta.

Deployment stage [New Stage] ▾

Stage name* test

Stage description Testing things

Deployment description

Cancel Deploy

And get the url in from the "Invoke URL" field. For us, it's <https://krs5gp9clb.execute-api.eu-west-1.amazonaws.com/test>

Now we can test it from the terminal:

```
URL="https://krs5gp9clb.execute-api.eu-west-1.amazonaws.com/test"
REQUEST='{
  "url":
    "https://raw.githubusercontent.com/alexeygrigorev/clothing-dataset-small/master/test/pants/4aabd82c-82e1-4181-a84d-d0c6e550d26d.jpg"
}'
```

```
curl -X POST \
  -H "Content-Type: application/json" \
  --data "${REQUEST}" \
  "${URL}/predict | jq
```

The response:

```
{
  "dress": -1.8682900667190552,
  "hat": -4.7612457275390625,
  "longsleeve": -2.3169822692871094,
  "outwear": -1.062570571899414,
  "pants": 9.88715648651123,
```

```
"shirt": -2.8124303817749023,  
"shoes": -3.66628360748291,  
"shorts": 3.2003610134124756,  
"skirt": -2.6023387908935547,  
"t-shirt": -4.835044860839844  
}
```

Now it's working!